

Random Forest Algorithm Intuition:

Random Forest is an ensemble learning method that builds multiple decision trees during training and merges their predictions to improve accuracy and prevent overfitting. Each decision tree is built on a random subset of features and data samples, and the final prediction is based on the combined predictions of all trees.

Advantages of Random Forest:

High Accuracy: Random Forests generally provide high accuracy due to the combination of multiple decision trees.

Handles Missing Values: Random Forests can handle missing data without imputation.

Reduced Overfitting: The randomness in feature selection and data sampling reduces overfitting compared to individual decision trees.

Feature Importance: It can provide insights into feature importance, aiding in feature selection.

Disadvantages of Random Forest:

Complexity: Random Forests can be computationally intensive and may require tuning of hyperparameters.

Less Interpretability: While they offer feature importance, interpreting individual tree decisions can be challenging.

Training Time: Training a Random Forest model can take longer compared to simpler algorithms like Decision Trees.

Difference between Random Forests and Decision Trees:

Random Forests are an ensemble of decision trees, while Decision Trees are single trees.

Random Forests introduce randomness in feature selection and data sampling to reduce overfitting, whereas Decision Trees can easily overfit on training data.

Relationship between Random Forests and Nearest Neighbors:

Random Forests and k-Nearest Neighbors (KNN) are different algorithms.

Random Forests use decision trees for classification, while KNN classifies based on the majority class among its k-nearest neighbors.

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, accu
from sklearn.preprocessing import StandardScaler
```

import and analyse data

```
In [ ]: dataset = pd.read_csv('heart-disease.csv')
```

```
In [ ]: dataset.head()
```

```
Out[ ]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	63	1	3	145	233	1	0	150	0	2.3	0	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	

```
In [ ]: dataset.isnull().sum()
```

```
Out[ ]: age          0
sex            0
cp             0
trestbps       0
chol           0
fbs            0
restecg        0
thalach        0
exang          0
oldpeak        0
slope          0
ca             0
thal           0
target         0
dtype: int64
```

```
In [ ]: dataset.describe()
```

```
Out[ ]:
```

	age	sex	cp	trestbps	chol	fbs	target
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.516832
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.500000
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000

since dataset is clean and preprocessed correctly already we'll move to next steps

feature engineering

```
In [ ]: numerical_features = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']
        scaler = StandardScaler()
        # scale numerical values to comparable ranges
        dataset[numerical_features] = scaler.fit_transform(dataset[numerical_feat
```

```
In [ ]: dataset.head()
```

```
Out[ ]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
0	0.952197	1	3	0.763956	-0.256334	1	0	0.015443	0	1.087338
1	-1.915313	1	2	-0.092738	0.072199	0	1	1.633471	0	2.122573
2	-1.474158	0	1	-0.092738	-0.816773	0	0	0.977514	0	0.310912
3	0.180175	1	1	-0.663867	-0.198357	0	1	1.239897	0	-0.206705
4	0.290464	0	0	-0.663867	2.082050	0	1	0.583939	1	-0.379244


```
In [ ]: X = dataset.drop(['target'], axis=1)
        y = dataset['target']
```

```
In [ ]: Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, test_size=0.2, rand
```

train using default params

```
In [ ]: def_model = RandomForestClassifier()
```

```
In [ ]: def_model.fit(Xtrain, ytrain)
        y_pred1 = def_model.predict(Xtest)
        cm1 = confusion_matrix(ytest, y_pred1)
        disp1 = ConfusionMatrixDisplay.from_predictions(ytest, y_pred1)
```



```
In [ ]: print(classification_report(ytest, y_pred1))
```

	precision	recall	f1-score	support
0	0.85	0.85	0.85	27
1	0.88	0.88	0.88	34
accuracy			0.87	61
macro avg	0.87	0.87	0.87	61
weighted avg	0.87	0.87	0.87	61

```
In [ ]: accuracy_score(ytest, y_pred1)
```

```
Out[ ]: 0.8688524590163934
```

Precision

For class 0 (no heart disease), the precision is 0.78, meaning that when the model predicts no heart disease, it is correct 78% of the time.

For class 1 (presence of heart disease), the precision is 0.82, indicating that when the model predicts heart disease, it is correct 82% of the time.

Recall

The recall for class 0 is also 0.78, which means that of all the actual no heart disease cases, the model correctly identifies 78%.

The recall for class 1 is 0.82, so the model correctly

identifies 82% of all actual heart disease cases.

F1-Score

The F1-score for both classes is quite balanced, with 0.78 for class 0 and 0.82 for class 1. This score gives a combined idea of how precise and robust the model's predictions are.

Accuracy

The overall accuracy of the model is 0.80, meaning the model correctly predicts the presence or absence of heart disease 80% of the time across all cases.


Analysis

this model shows a good balance between precision and recall, indicating it is equally skilled at identifying positive cases among the actual positives and minimizing false positives.
The balanced macro and weighted averages indicate that the model performs consistently across both classes

now using different params

```
In [ ]: model = RandomForestClassifier(n_estimators=500, n_jobs=-1, random_state=0)
        model.fit(Xtrain, ytrain)
        y_pred2 = model.predict(Xtest)
        cm = confusion_matrix(ytest, y_pred2)
        disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=np.unique(ytest))

In [ ]: disp.plot()
        plt.show()
```



```
In [ ]: print(classification_report(ytest, y_pred2))
```

	precision	recall	f1-score	support
0	0.85	0.81	0.83	27
1	0.86	0.88	0.87	34
accuracy			0.85	61
macro avg	0.85	0.85	0.85	61
weighted avg	0.85	0.85	0.85	61

Precision

Class 0 (No Heart Disease): Precision has improved to 0.85, meaning the model is more accurate in predicting negative cases

Class 1 (Presence of Heart Disease): Precision is slightly higher at 0.86, indicating a slight improvement in correctly predicting positive cases.

Recall

Class 0: Recall has improved to 0.81, suggesting that the model is now better at identifying all actual negative cases as negative.

Class 1: Recall has increased to 0.88, meaning the model has become more effective at catching true positive cases.

F1-Score

The F1-scores have increased to 0.83 for class 0 and 0.87 for class 1, reflecting a balanced improvement in precision and recall for both classes

Accuracy

The overall accuracy has risen to 0.85. This improvement suggests that the adjustments to the number of trees and maximum depth have made the model better at classifying both conditions correctly across all predictions.

In []: